

**Synopsis
Of
Transformer Architecture**

**Submitted for NUES (IT 669) Term
Paper 1**

M.C.A. (Software Engineering)

By

**Barun Wason
Application No. 105242000356**

Under the supervision of
**Dr. Virendra P.
Vishwakarma
(Professor, USICT)**

University School of Information Communication and Technology



**Guru Gobind Singh Indraprastha University
Dwarka, Sector 16C, New Delhi**

INTRODUCTION

Transformer Architecture: Revolutionizing Neural Networks in Natural Language Processing

The Transformer architecture, introduced in the groundbreaking paper “*Attention is All You Need*” by Vaswani et al. (2017), has become the foundation for modern large language models (LLMs) like GPT, BERT, and more. This architecture addresses the limitations of traditional recurrent neural networks (RNNs) by eliminating sequential processing and focusing on parallelization, scalability, and enhanced learning capabilities. This paper explores the inner workings of the Transformer model and its pivotal components: **positional encoding**, **attention mechanisms**, and **self-attention**.

Objective:

The goal of this term paper is to:

1. Provide a comprehensive overview of the Transformer architecture.
2. Compare its effectiveness with traditional models like RNNs and LSTMs.
3. Explain how the introduction of **self-attention** and **positional encoding** optimizes information processing in neural networks, making Transformers more efficient in handling long-range dependencies.

Key Components:

- **Positional Encoding:** Since Transformers do not have a sequential structure like RNNs, positional encodings are added to input embeddings to inject information about the position of words in a sentence. This ensures that the model can discern the order of tokens.
- **Attention Mechanism:** The attention mechanism allows the model to focus on important parts of the input by weighing different tokens' importance. It enables the model to handle variable-length inputs more efficiently, unlike RNNs which suffer from long-term dependency issues.
- **Self-Attention:** Self-attention enables the model to relate different positions of the input sequence to each other and determine which tokens are most relevant to the current token, enhancing parallelization and contextual understanding.

Comparison with RNNs:

Transformers outperform RNNs and LSTMs in several ways:

- **Parallelization:** Unlike RNNs, which process sequences step-by-step, Transformers process input tokens simultaneously, significantly reducing training time.
- **Handling Long Sequences:** RNNs struggle with vanishing gradient problems in long sequences, while Transformers, through attention mechanisms, excel in modeling long-range dependencies.
- **Scalability:** Transformers can easily scale to much larger models like GPT-3, whereas RNNs face challenges in scaling due to their sequential nature.

Conclusion:

The Transformer architecture has dramatically changed the landscape of NLP and neural networks by addressing the limitations of previous models. Its ability to model complex relationships and scale to large datasets has made it the backbone of modern AI applications, from language translation to large language models.

References:

- Vaswani, A., et al. (2017). *Attention is All You Need*. arXiv:1706.03762.

